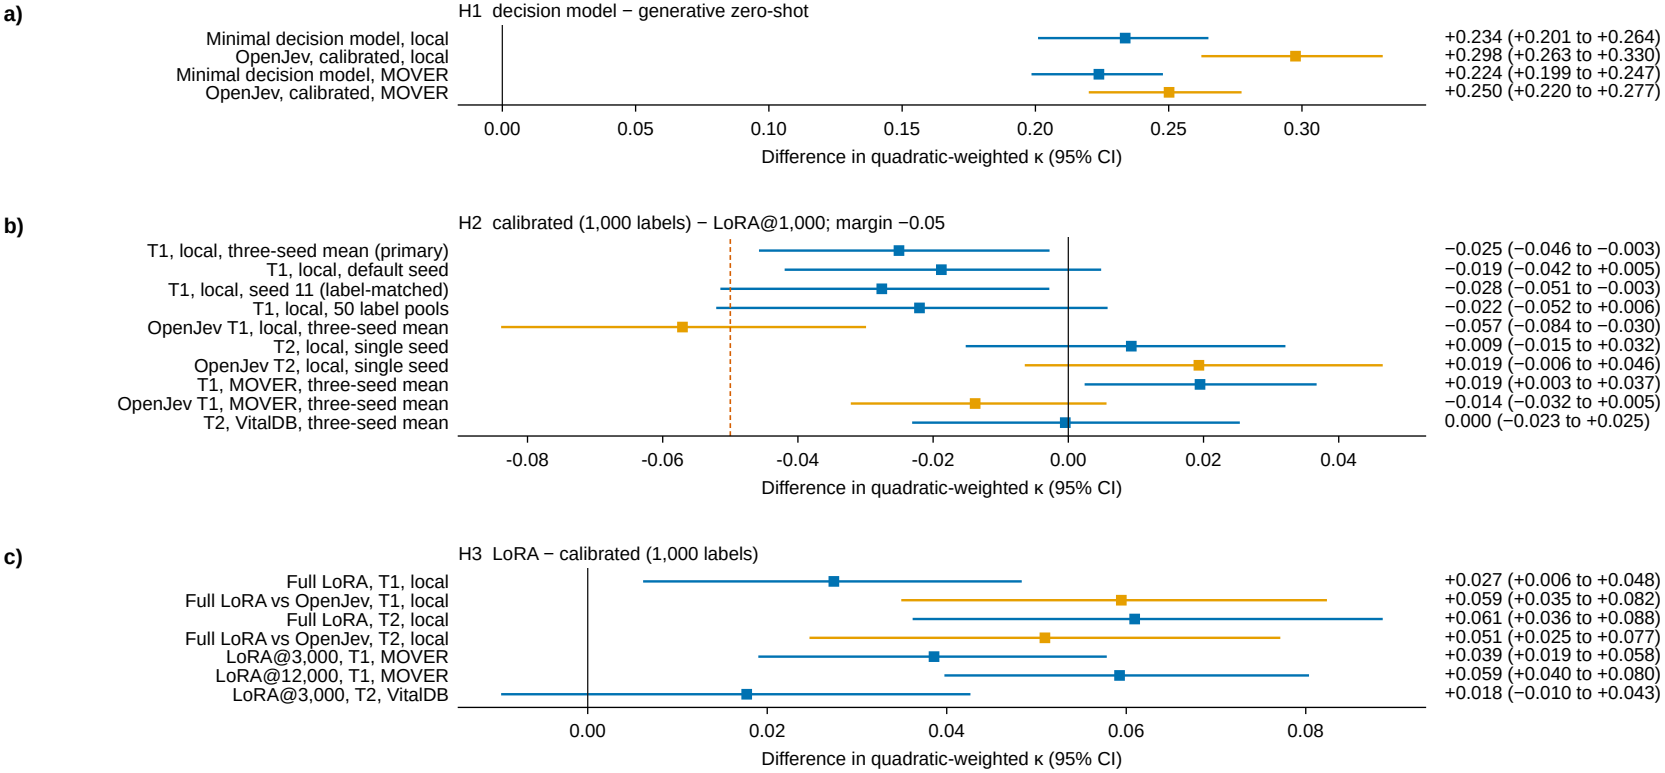
a)
H1 decision model − generative zero-shot
Minimal decision model, local
OpenJev, calibrated, local
Minimal decision model, MOVER
OpenJev, calibrated, MOVER
+0.234 (+0.201 to +0.264)
+0.298 (+0.263 to +0.330)
+0.224 (+0.199 to +0.247)
+0.250 (+0.220 to +0.277)
0.00
0.05
0.10
0.15
0.20
0.25
0.30
Difference in quadratic-weighted κ (95% CI)
b)
H2 calibrated (1,000 labels) − LoRA@1,000; margin −0.05
T1, local, three-seed mean (primary)
T1, local, default seed
T1, local, seed 11 (label-matched)
T1, local, 50 label pools
OpenJev T1, local, three-seed mean
T2, local, single seed
OpenJev T2, local, single seed
T1, MOVER, three-seed mean
OpenJev T1, MOVER, three-seed mean
T2, VitalDB, three-seed mean
−0.025 (−0.046 to −0.003)
−0.019 (−0.042 to +0.005)
−0.028 (−0.051 to −0.003)
−0.022 (−0.052 to +0.006)
−0.057 (−0.084 to −0.030)
+0.009 (−0.015 to +0.032)
+0.019 (−0.006 to +0.046)
+0.019 (+0.003 to +0.037)
−0.014 (−0.032 to +0.005)
0.000 (−0.023 to +0.025)
-0.08
-0.06
-0.04
-0.02
0.00
0.02
0.04
Difference in quadratic-weighted κ (95% CI)
c)
H3 LoRA − calibrated (1,000 labels)
Full LoRA, T1, local
Full LoRA vs OpenJev, T1, local
Full LoRA, T2, local
Full LoRA vs OpenJev, T2, local
LoRA@3,000, T1, MOVER
LoRA@12,000, T1, MOVER
LoRA@3,000, T2, VitalDB
+0.027 (+0.006 to +0.048)
+0.059 (+0.035 to +0.082)
+0.061 (+0.036 to +0.088)
+0.051 (+0.025 to +0.077)
+0.039 (+0.019 to +0.058)
+0.059 (+0.040 to +0.080)
+0.018 (−0.010 to +0.043)
0.00
0.02
0.04
0.06
0.08
Difference in quadratic-weighted κ (95% CI)